

Homework 1

Shun-Lung Chang, Dilip Hiremath

1. Before you address your real task you just want to get a bit acquainted with the data. Read the data into R and use the function `summary` to get a first overview on the data.

```
load("data/Forbes2016.Rdata")
summary(Forbes2016)
```

Rank		Company		Country	
Min.	: 1	Merck	: 2	United States	:586
1st Qu.:	500	3i Group	: 1	Japan	:219
Median	:1000	3M	: 1	China	:200
Mean	:1001	77 Bank	: 1	United Kingdom	: 92
3rd Qu.:	:1501	AAC Technologies Holdings	: 1	South Korea	: 67
Max.	:2000	Aareal Bank	: 1	France	: 61
		(Other)	:1994	(Other)	:776

Sales		Profits		Assets		Market.Value	
Min.	: 0.24	Min.	:-23.100	Min.	: 0.98	Min.	: 0.44
1st Qu.:	4.30	1st Qu.:	1.500	1st Qu.:	10.40	1st Qu.:	6.30
Median	: 8.90	Median	: 3.380	Median	: 21.70	Median	: 10.90
Mean	: 17.84	Mean	: 3.562	Mean	: 80.81	Mean	: 22.35
3rd Qu.:	17.90	3rd Qu.:	5.570	3rd Qu.:	49.15	3rd Qu.:	21.90
Max.	:482.10	Max.	: 53.700	Max.	:3420.30	Max.	:586.00
NA's	:4			NA's	:1		

(a) How many companies in the data set are located in China?

There are 200 companies located in China, as one can see from the summary table.

(b) How many companies in the data set have missing values (i.e. are listed as NA)? For which variables are data missing?

If we look at *Assets* and *Market.Value* variables, there is 1 missing value in both variables. The row containing missing values is shown below.

```
kable(Forbes2016[is.na(Forbes2016$Assets), ])
```

	Rank	Company	Country	Sales	Profits	Assets	Market.Value
1840	1840	Banque nationale de Belgique	Belgium	1.7	9.03	NA	1.3

```
kable(Forbes2016[is.na(Forbes2016$Market.Value), ])
```

	Rank	Company	Country	Sales	Profits	Assets	Market.Value
685	685	Gree Electric Appliances	China	19.7	2.3	26.8	NA

In addition, there exist four rows containing missing values in *Sales* variable. These rows are shown below.

```
kable(Forbes2016[is.na(Forbes2016$Sales), ])
```

	Rank	Company	Country	Sales	Profits	Assets	Market.Value
672	672	Porsche Automobil Holding	Germany	NA	3.50	34.9	16.60
1072	1072	Kyushu Financial Group	Japan	NA	8.12	73.0	2.60
1398	1398	RHB Capital	Malaysia	NA	3.87	53.7	4.90
1988	1988	Ivanhoe Mines	Canada	NA	7.28	1.0	6.34

(c) Which four countries have the top most companies in the Forbes' list?

The first top countries are United States, Japan, China, and United Kingdom.

(d) Which share of companies in the data set have a market value smaller than or equal to 6.30 billion USD?

The share is 25.24%, and the value can be counted in the following code.

```
sum(Forbes2016$Market.Value <= 6.3, na.rm = TRUE) / nrow(Forbes2016)
```

```
[1] 0.2523738
```

(e) Which share of companies in the data set have assets that are larger than or equal to 10.40 billion USD?

The share is 75.01%, and the value can be counted in the following code.

```
sum(Forbes2016$Assets >= 10.4, na.rm = TRUE) / nrow(Forbes2016)
```

```
[1] 0.7501249
```

2. Let us have a closer look at the variable types involved here. Looking at the content of each variable:

(a) Which of the variables in the data set are discrete?

The discrete variables are *Rank*, *Company*, and *Country*.

(b) Which of the variables are continuous?

The discrete variables are *Sales*, *Profits*, *Assets*, and *Market.Value*.

(c) Which discrete ones can be ordered?

The *Rank* variable can be ordered.

3. Which types are used in R to store the variables in the data frame?

The following code returns the class of each variable.


```
sapply(Forbes2016, class)
```

Rank	Company	Country	Sales	Profits
"integer"	"factor"	"factor"	"numeric"	"numeric"
Assets	Market.Value			
"numeric"	"numeric"			

4. Look at the summary statistics to see how the variable class affects the output. Which variable classes produce the same summary output? Which produce a different one?

5. Draw a histogram for the variable Profit using the default number of bins and color the bins in 'springgreen'. Add a meaningful title to the plot.

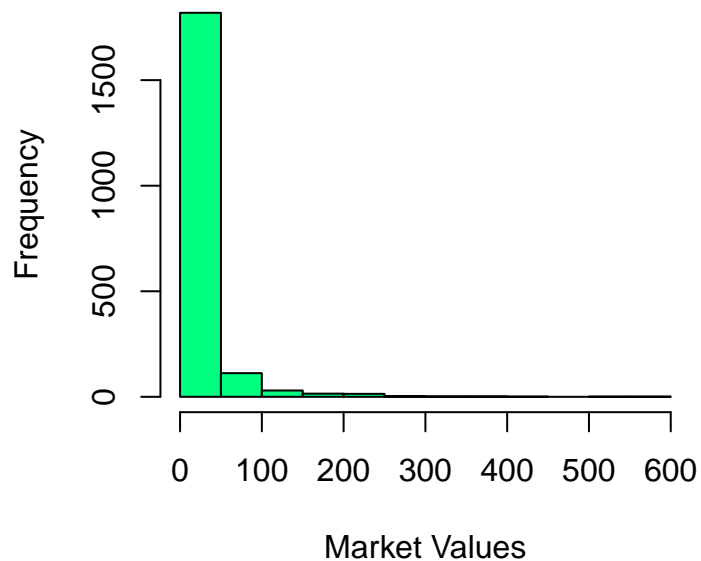
```
hist(Forbes2016$Profits, col = "springgreen",  
     main = "The Count of Profits", xlab = "Profits")
```



(a) Is the distribution of market values in the data set symmetric?

```
hist(Forbes2016$Market.Value, col = "springgreen",  
     main = "The Count of Market Values", xlab = "Market Values")
```

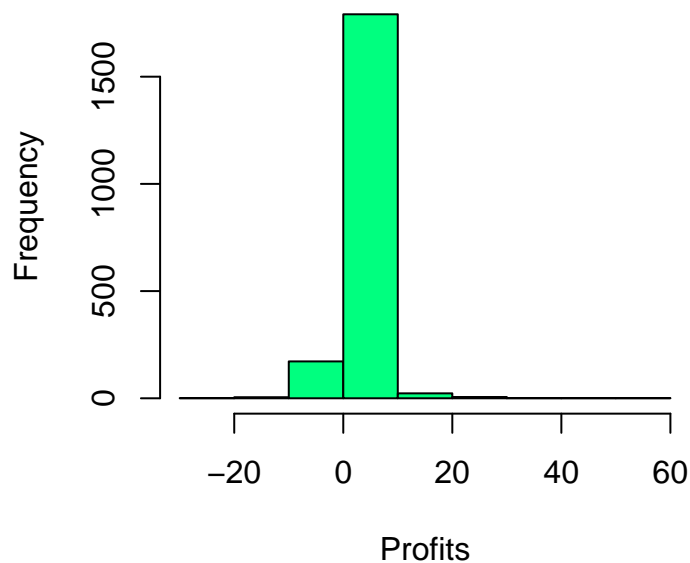
The Count of Market Values



(b) Does the selected number of bins yield a good graphical representation of the data? Try out 7, 50 and 100 bins! Is any of these three options more suited?


```
hist(Forbes2016$Profits, col = "springgreen", breaks = 7,  
     main = "The Count of Profits (bins = 7)", xlab = "Profits")
```

The Count of Profits (bins = 7)




```
hist(Forbes2016$Profits, col = "springgreen", breaks = 50,  
     main = "The Count of Profits (bins = 50)", xlab = "Profits")
```

The Count of Profits (bins = 50)



```
hist(Forbes2016$Profits, col = "springgreen", breaks = 100,  
     main = "The Count of Profits (bins = 100)", xlab = "Profits")
```


The Count of Profits (bins = 100)



6. Draw a boxplot for the variable Profit.

```
boxplot(Forbes2016$Profits, horizontal = TRUE, main = "Boxplot of Profits")
```

Boxplot of Profits



- Compare the boxplot with the histograms above.
- Which graphical representation of the data do you find more helpful?
- How many companies do have profits above 40 billion USD? Can you find out which companies these are?

```
Forbes2016[Forbes2016$Profits >= 40, ]$Company
```

```
[1] ICBC Apple  
2000 Levels: 3i Group 3M 77 Bank AAC Technologies Holdings ... Zurich Insurance Group
```

7. You have a closer look at profits of the companies.

- First, you want to know which profit is needed such that a company belongs to the top two percent.

```
quantile(Forbes2016$Profits, probs = 0.98)
```

```
98%  
9.87
```

- Which percentage of companies has negative profits?

```
sum(Forbes2016$Profits < 0) / nrow(Forbes2016)
```

```
[1] 0.08845577
```


(c) How large is the interquartile range in sales for the companies on Forbes' list?

```
IQR(Forbes2016$Sales, na.rm = TRUE)
```

```
[1] 13.6
```

8. Now you start turning towards your real task. Create a subset of the data set comprising all companies that have their legal residence in one of the following countries: Brazil, India, South Korea, France, Germany or Mexico. Draw a boxplot for the variable Profits splitted by country for these companies.

```
sub_forbes2016 <- subset(Forbes2016,  
                          Country %in% c("Brazil", "India", "South Korea", "Franc", "Germany", "Mexico"))  
boxplot(Profits ~ as.character(Country), data = sub_forbes2016, horizontal = TRUE)
```



(a) How much difference in medians between countries is shown in the plot?

(b) Which country shows the largest median profit?

The plot indicates that the largest median profit lies in Brazil.

(c) Which country/countries show the most symmetric distributions for profit?

South Korea shows the most symmetric distributions for profit in the plot.

(d) Create a frequency table for the variable Country sorted in decreasing order. Which three countries – among those selected – have the top most companies in Forbes' list?

```
sort(table(as.character(sub_forbes2016$Country)), decreasing = TRUE)
```

South Korea	India	Germany	Brazil	Mexico
67	56	50	19	15

9. Now you start with your real task for the subset data created in Question 8.

(a) What is the median profit of the selected companies?

```
median(sub_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 2.95
```

(b) What is the average profit of the selected companies?

```
mean(sub_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 3.037536
```

(c) What is the average distance for a company's profit to the average profit (i.e. what is the standard deviation of profits)?

```
sd(sub_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 3.81334
```

(d) Does the top third of selected companies have profits that are at least as high as the median profit for all companies on the Forbes list?

```
aggregate(sub_forbes2016$Profits, list(sub_forbes2016$Country), median)
```

```

      Group.1      x
1      Brazil 3.770
2     Germany 3.450
3       India 2.805
4       Mexico 2.380
5 South Korea 2.700

```

10. In order to fulfil your task you finally draw a random sample of size 80 from the list of your selected companies. Create the resulting data frame!

```
set.seed(42)
sampled_forbes2016 <- sub_forbes2016[sample(1:nrow(sub_forbes2016), 80), ]
```

(a) Compute the mean profit for your sample?

```
mean(sampled_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 2.512125
```


(b) Compute the median profit for your sample?

```
median(sampled_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 2.25
```

(c) Compute the minimum profit for your sample?

```
min(sampled_forbes2016$Profits, na.rm = TRUE)
```

```
[1] -13.2
```

(d) Compute the maximum profit for your sample?

```
max(sampled_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 16.5
```

(e) Compute the 90%-quantile for your sample?

```
quantile(sampled_forbes2016$Profits, probs = 0.9)
```

```
90%  
7.335
```

11. Are median and mean profits of your sample below or above the median and mean market value for all companies in the selected countries?

12. Compute the standard deviation of profits for your sample. Is it smaller or larger than the standard deviation in profits for all selected companies?


```
sd(sampled_forbes2016$Profits, na.rm = TRUE)
```

```
[1] 4.201943
```

13. Draw a boxplot, a histogram, and a density plot of the profit for your sample. Can you draw any substantial information from your plot? (e.g about skewness of the distribution, outliers, quartiles)


```
boxplot(sampled_forbes2016$Profits, horizontal = TRUE,  
        main = "Boxplot of Profits for Sampled Data")
```

Boxplot of Profits for Sampled Data



```
hist(sampled_forbes2016$Profits,  
     main = "Histogram of Profits for Sampled Data", xlab = "Profits")
```

Histogram of Profits for Sampled Data



```
d <- density(sampled_forbes2016$Profits)  
plot(d, main = "Density of Profits for Sampled Data", xlab = "Profits")
```

Density of Profits for Sampled Data

